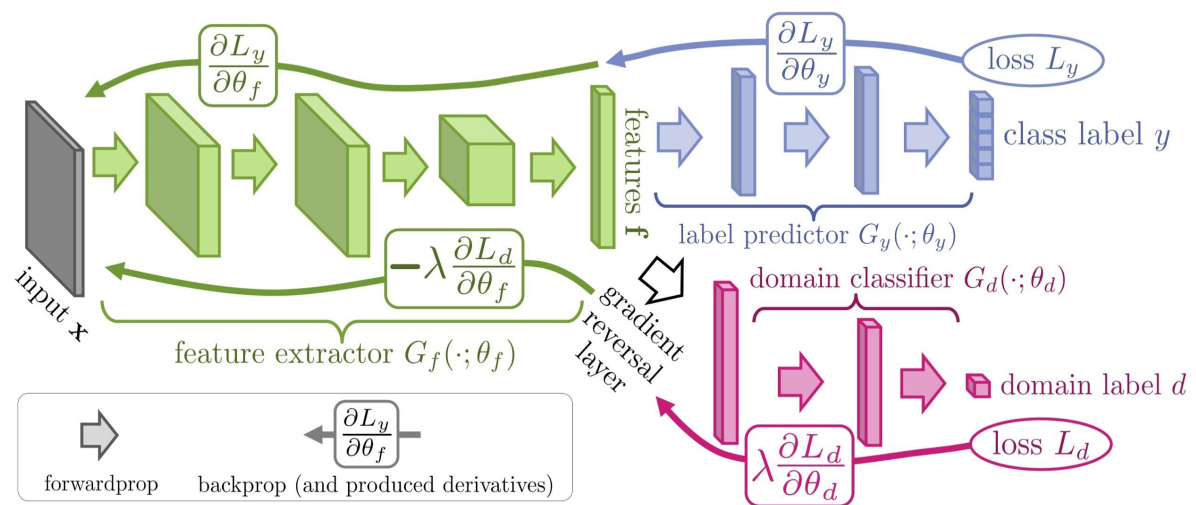


You can find my results on my github: <https://github.com/Barrel-Titor/homework-MLDL>

Implementation of DANN

There are 3 parts of DANN: feature extractor G_f , label predictor G_y and domain classifier G_d

The idea is like training GAN, but according to the paper, a gradient reversal layer is used for combining feature extractor G_f and domain classifier G_d .



The implementation of DANN is as follows:

```
1 from PACS.gradient_reversal_example import ReverseLayerF
2 from torchvision.models import AlexNet
3 from torchvision.models.utils import load_state_dict_from_url
4 from copy import deepcopy
5
6 model_urls = {
7     'alexnet': 'https://download.pytorch.org/models/alexnet-owt-
8     4df8aa71.pth',
9 }
10
11 class DANN(AlexNet):
12     def __init__(self):
13         super(DANN, self).__init__()
14         self.domain_clf = nn.Sequential(
15             nn.Dropout(),
16             nn.Linear(256 * 6 * 6, 4096),
17             nn.ReLU(inplace=True),
18             nn.Dropout(),
19             nn.Linear(4096, 4096),
20             nn.ReLU(inplace=True),
21             nn.Linear(4096, 2),
22         )
23
```

```

24     def forward(self, x, alpha=None):
25         x = self.features(x)
26         x = self.avgpool(x)
27         # Flatten the features:
28         x = x.view(x.size(0), -1)
29         # If we pass alpha, we can assume we are training the discriminator
30         if alpha is not None:
31             # gradient reversal layer (backward gradients will be reversed)
32             reverse_feature = ReverseLayerF.apply(x, alpha)
33             discriminator_output = self.domain_clf(x)
34             return discriminator_output
35         # If we don't pass alpha, we assume we are training with
supervision
36         else:
37             # do something else
38             class_outputs = self.classifier(x)
39             return class_outputs
40
41     def dann(pretrained=False, progress=True, **kwargs):
42         r"""AlexNet model architecture from the
43         `One weird trick...` <https://arxiv.org/abs/1404.5997>`_ paper.
44         Args:
45             pretrained (bool): If True, returns a model pre-trained on ImageNet
46             progress (bool): If True, displays a progress bar of the download
to stderr
47         """
48         model = DANN(**kwargs)
49         if pretrained:
50             state_dict = load_state_dict_from_url(model_urls['alexnet'],
51                                                    progress=progress)
52             model.load_state_dict(state_dict, strict=False)
53         return model

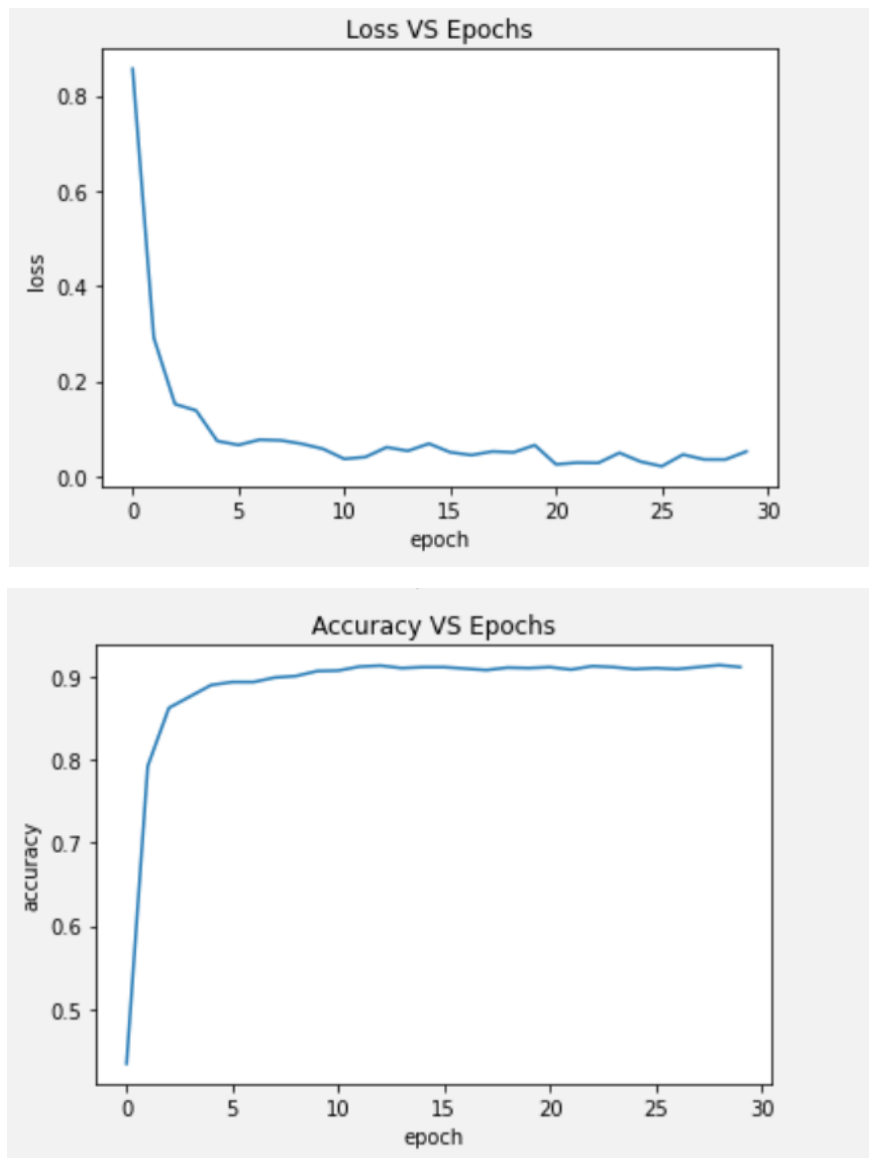
```

When we want data flow from feature extractor G_f to label predictor G_y , we simply use `outputs = net(input)`

When we want data flow from feature extractor G_f to domain classifier G_d , we additionally transfer a hyperparameter `ALPHA` which should be optimized later: `outputs = net.forward(input, ALPHA)`

Train on P and test on A without adaptation

Training without domain adaptation is just like Homework2, using AlexNet.



1 | `test_accuracy` # 0.46484375

Train on P and test on A with DANN

Each training epoch is divided into 3 parts before calling `optimizer.step()`

1. Data flow from G_f to G_y , training G_y by source data and labels;
2. Data flow from G_f to G_d , training G_d by **source data and labels of all 0**, and with the help of gradient reversal layer, training G_f to cheat the domain classifier G_d
3. Data flow from G_f to G_d , training G_d by **target data and labels of all 1**, and with the help of gradient reversal layer, training G_f to cheat the domain classifier G_d

Before training, there is an additional step, i.e. copying weights of label classifier into domain classifier

```
1 net.domain_clf[1].weight.data = net.classifier[1].weight.data
2 net.domain_clf[1].bias.data = net.classifier[1].bias.data
3
4 net.domain_clf[4].weight.data = net.classifier[4].weight.data
5 net.domain_clf[4].bias.data = net.classifier[4].bias.data
```

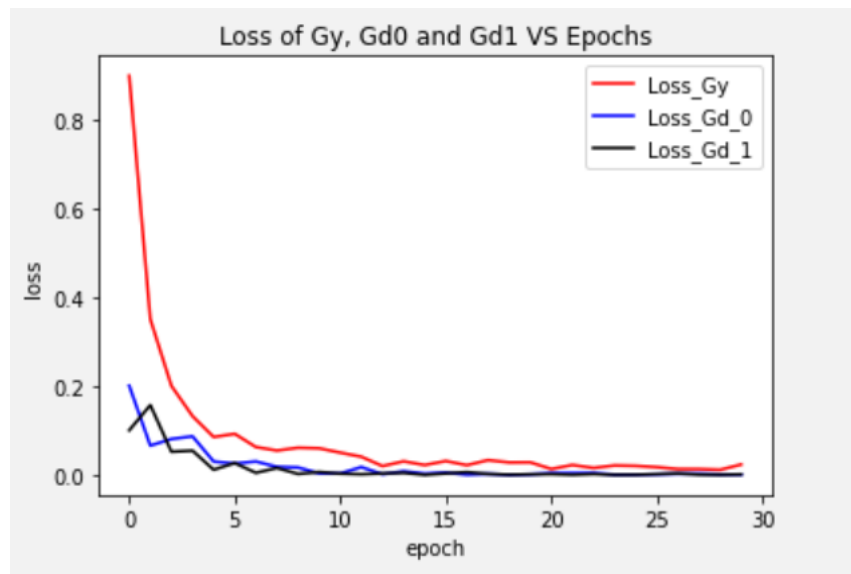
Training phase is as follows:

```
1  # By default, everything is loaded to cpu
2  net = net.to(DEVICE) # this will bring the network to GPU if DEVICE is cuda
3
4  cudnn.benchmark # Calling this optimizes runtime
5
6  loss_Gy_hist = []
7  loss_Gd_0_hist = []
8  loss_Gd_1_hist = []
9
10 # Start iterating over the epochs
11 for epoch in range(NUM_EPOCHS):
12     print('Starting epoch {}/{}'.format(epoch+1, NUM_EPOCHS),
13           scheduler.get_last_lr())
14
15     # Iterate over the dataset
16     for images, labels in train_data_loader:
17         # Bring data over the device of choice
18         images = images.to(DEVICE)
19         labels = labels.to(DEVICE)
20
21         net.train() # Sets module in training mode
22
23         # PyTorch, by default, accumulates gradients after each backward pass
24         # We need to manually set the gradients to zero before starting a new
25         # iteration
26         optimizer.zero_grad() # Zero-ing the gradients
27
28         # 3B.1 Train Gy
29         # Forward pass to the network
30         outputs = net(images)
31
32         # Compute loss based on output and ground truth
33         loss = criterion(outputs, labels)
34
35         loss.backward() # backward pass: computes gradients
36
37         # 3B.2 Train Gd by forwarding source data
38         label_outputs = net.forward(images, ALPHA)
39         targets = torch.zeros(labels.shape, dtype=int).to(DEVICE)
40
41         loss_Gd_0 = criterion(label_outputs, targets)
42         loss_Gd_0.backward()
43
44         # 3B.3 Train Gd by forwarding target data
45         test_images, test_labels = next(iter(test_data_loader))
46         test_images = test_images.to(DEVICE)
47         test_labels = test_labels.to(DEVICE)
48
49         test_label_outputs = net.forward(test_images, ALPHA)
50         targets = torch.ones(test_labels.shape, dtype=int).to(DEVICE)
51
52         loss_Gd_1 = criterion(test_label_outputs, targets)
53         loss_Gd_1.backward()
54
55         # update weights based on accumulated gradients
```

```

54     optimizer.step()
55
56     # Step the scheduler
57     scheduler.step()
58
59     # Log loss
60     print('Gy Loss {}'.format(loss.item()))
61     print('Gd0 Loss {}'.format(loss_Gd_0.item()))
62     print('Gd1 Loss {}'.format(loss_Gd_1.item()))
63
64     # Record loss and accuracy after each epoch
65     loss_Gy_hist.append(loss.item())
66     loss_Gd_0_hist.append(loss_Gd_0)
67     loss_Gd_1_hist.append(loss_Gd_1)

```



```

1 | test_accuracy    # 0.49169921875

```

Hyperparameter optimization

Since there are too many outputs during training, I use logs to record the prints.

```

1 | logger = logging.getLogger(__name__)
2 | logger.setLevel(logging.DEBUG)
3 | logname = 'DANN.log'
4 | handler = logging.FileHandler(logname)
5 | formatter = logging.Formatter(
6 |     '%(asctime)s - %(levelname)s - %(lineno)d - %(message)s'
7 | )
8 | handler.setFormatter(formatter)
9 | logger.addHandler(handler)
10
11 | logger.info('Starting epoch {}/{}, LR = {}, alpha = {}'.format(
12 |     epoch+1, NUM_EPOCHS, scheduler.get_last_lr(), alpha
13 | ))
14 | ...


```

I have

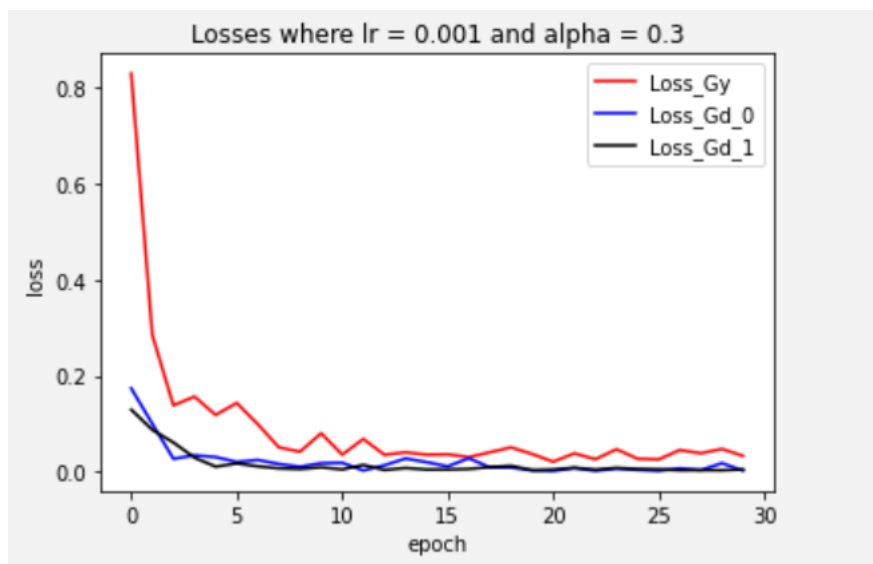
```
1 lr_choices = [1e-3, 1e-4, 1e-5]
2 alpha_choices = [0.1, 0.3, 1, 3, 10]
```

and perform the grid search.

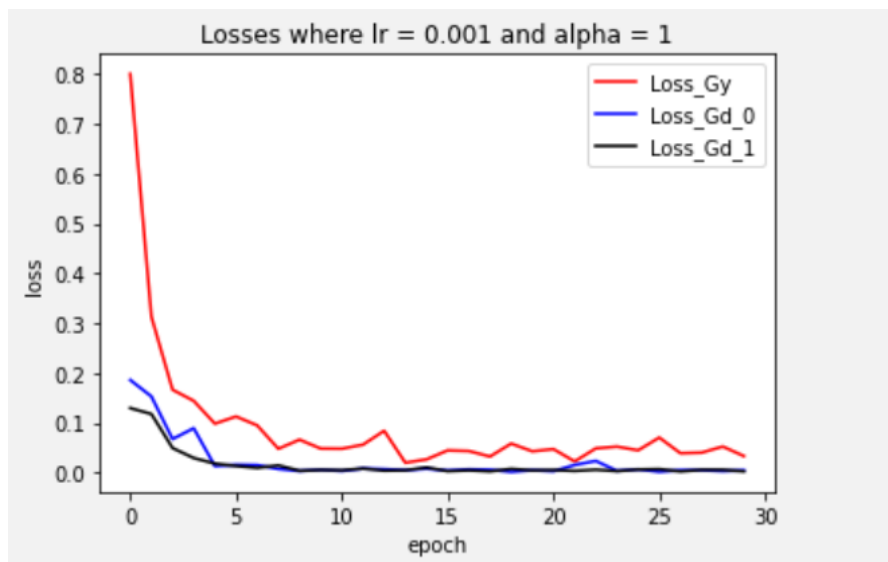
```
1 for lr in lr_choices:
2     for alpha in alpha_choices:
3         # prepare network and optimizer
4         ...
5
6         # training, with three parts as mentioned before
7         ...
8
9         # testing
10        ...
11
12        # save the lr, alpha and net parameters with best accuracy
13        ...
```




```
1 test_accuracy # 0.482421875
```




1 | test_accuracy # 0.48291015625




1 | test_accuracy # 0.45263671875



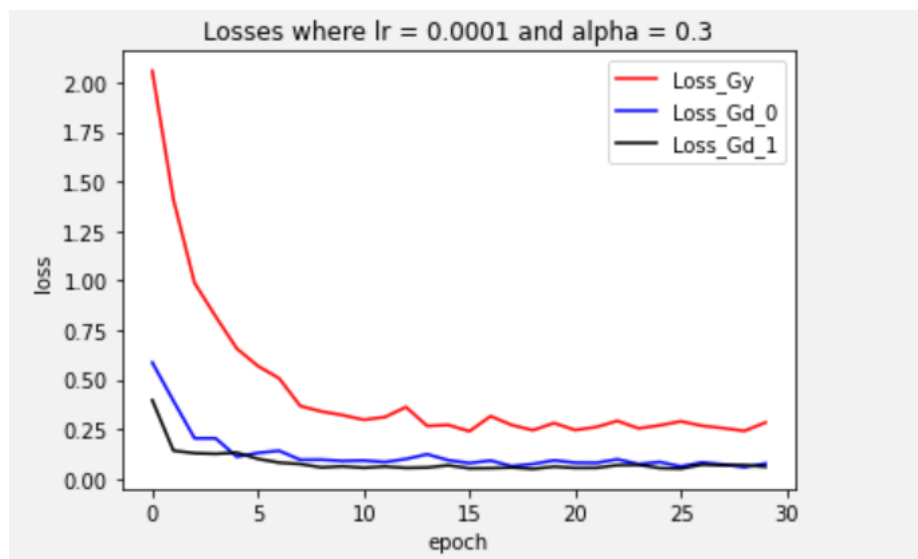
1 | test_accuracy # 0.48388671875




1 | test_accuracy # 0.484375




1 | test_accuracy # 0.45947265625




1 | test_accuracy # 0.42431640625




1 | test_accuracy # 0.43505859375




1 | test_accuracy # 0.3798828125




1 | test_accuracy # 0.43017578125



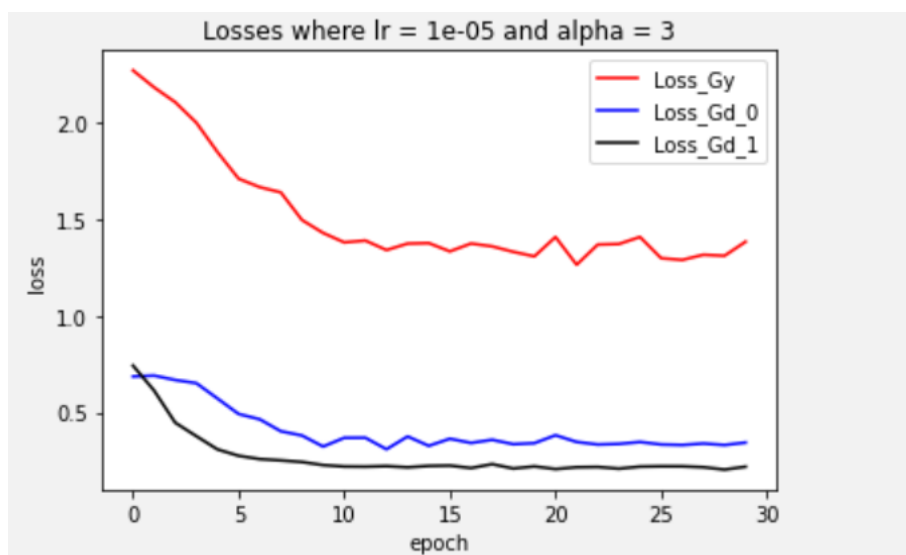
1 | test_accuracy # 0.25244140625



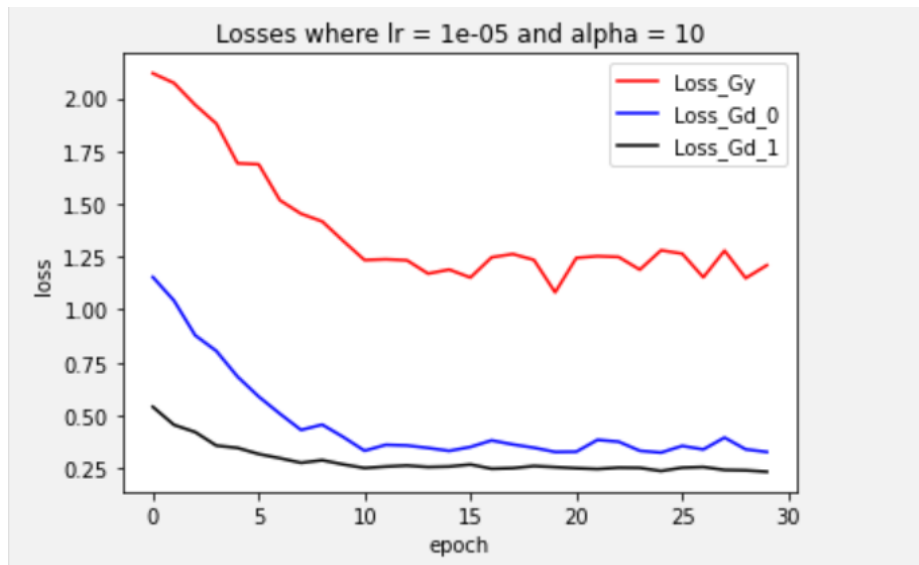
```
1 | test_accuracy # 0.228515625
```



```
1 | test_accuracy # 0.25537109375
```



```
1 | test_accuracy # 0.265625
```



```
1 | test_accuracy # 0.22802734375
```

We find that low `lr` behaves very bad. When `lr` is big enough, usually a big `alpha` performs a little bit better.

```
1 | best_lr, best_alpha, best_accuracy # (0.001, 10, 0.484375)
```